


Interval Scheduling

Interval scheduling.

- Job j starts at s_j and finishes at f_j .
- Two jobs **compatible** if they don't overlap.
- Goal: find maximum subset of mutually compatible jobs.



Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.



counterexample for earliest start time



counterexample for shortest interval



counterexample for fewest conflicts

Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.

- [Earliest start time] Consider jobs in ascending order of s_j .
- [Earliest finish time] Consider jobs in ascending order of f_j .
- [Shortest interval] Consider jobs in ascending order of $f_j - s_j$.
- [Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

Interval Scheduling: Greedy Algorithm

Greedy algorithm. Consider jobs in increasing order of finish time.
Take each job provided it's compatible with the ones already taken.

```

Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .
A ← ∅
set of jobs selected
for j = 1 to n {
    if (job j compatible with A)
        A ← A ∪ {j}
}
return A


```

Implementation. $O(n \log n)$.

Remember job j^* that was added last to A.

Job j is compatible with A if $s_j \geq f_{j^*}$.

Interval Scheduling



Interval Scheduling: Analysis

Theorem. Greedy algorithm is optimal.


Pf. (by contradiction)

Assume greedy is not optimal, and let's see what happens.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in the optimal solution with

$i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



Greedy algorithm stays ahead. Show that after each step of the greedy algorithm, its solution is at least as good as any other algorithm's.

Structural. Discover a simple "structural" bound asserting that every possible solution must have a certain value. Then show that your algorithm always achieves this bound.

Exchange argument. Gradually transform any solution to the one found by the greedy algorithm without hurting its quality.

Other greedy algorithms. Kruskal, Prim, Dijkstra, Huffman, ...

Greedy Algorithms

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.


Reverse-Delete algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Remark. All three algorithms produce an MST.


Cycle-Cut Intersection

Claim. A cycle and a cutset intersect in an even number of edges.



Cycle $C = 1-2, 2-3, 3-4, 4-5, 5-6, 6-1$
Cutset $D = 3-4, 3-5, 5-6, 5-7, 7-6$
Intersection $= 3-4, 5-6$

Pf. (by picture)



MST is fundamental problem with diverse applications.

Network design.

- telephone, electrical, hydraulic, TV cable, computer, road

Approximation algorithms for NP-hard problems.

- traveling salesperson problem, Steiner tree

Indirect applications.


- max bottleneck paths
- LDPC codes for error correction
- image registration with Renyi entropy
- learning salient features for real-time face verification
- reducing data storage in sequencing amino acids in a protein
- model locality of particle interactions in turbulent fluid flows
- autoconfig protocol for Ethernet bridging to avoid cycles in a network

Greedy Algorithms


Simplifying assumption. All edge costs c_e are distinct.

Cut property. Let S be any subset of nodes, and let e be the min cost edge with exactly one endpoint in S . Then the MST contains e .

Cycle property. Let C be any cycle, and let f be the max cost edge belonging to C . Then the MST does not contain f .



e is in the MST



f is not in the MST

Greedy Algorithms

Simplifying assumption. All edge costs c_e are distinct.

Cut property. Let S be any subset of nodes, and let e be the min cost edge with exactly one endpoint in S . Then the MST T^* contains e .

Pf. (exchange argument)

Suppose e does not belong to T^* , and let's see what happens.

Adding e to T^* creates a cycle C in T^* .


Edge e is both in the cycle C and in the cutset D corresponding to S

\Rightarrow there exists another edge, say f , that is in both C and D .

$T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree.

Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.


This is a contradiction. \bullet




Kruskal's Algorithm: Proof of Correctness

Kruskal's algorithm. [Kruskal, 1956]

- Consider edges in ascending order of weight.
- Case 1: If adding e to T creates a cycle, discard e according to cycle property.
- Case 2: Otherwise, insert $e = (u, v)$ into T according to cut property where S = set of nodes in u 's connected component.



Case 1



Case 2

Implementation: Kruskal's Algorithm

Implementation. Use the **union-find** data structure.


- Build set T of edges in the MST.
 - Maintain set for each connected component.
 - $O(m \log n)$ for sorting and $O(m \alpha(m, n))$ for union-find.
- $m \leq n^2 \Rightarrow \log m$ is $O(\log n)$ α is essentially a constant

```
Kruskal(G, c) {
  Sort edges weights so that  $c_1 \leq c_2 \leq \dots \leq c_m$ .
  T  $\leftarrow \emptyset$ 
  foreach (u  $\in$  V) make a set containing singleton u
  for i = 1 to m
    (u, v) =  $e_i$ 
    if (u and v are in different sets) {
      T  $\leftarrow T \cup \{e_i\}$ 
    }
  return T
```

Prim's Algorithm: Proof of Correctness

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize S = any node.
- Apply cut property to S .
- Add min cost edge in cutset corresponding to S to T , and add one new explored node u to S .



Implementation: Prim's Algorithm

Implementation. Use a priority queue ala Dijkstra.

- Maintain set of explored nodes S .
- For each unexplored node v , maintain attachment cost $a[v]$ = cost of cheapest edge v to a node in S .
- $O(n^2)$ with an array; $O(m \log n)$ with a binary heap.

```
Prim(G, c) {
  foreach (v  $\in$  V) a[v]  $\leftarrow \infty$ 
  Initialize an empty priority queue Q
  foreach (v  $\in$  V) insert v onto Q
  Initialize set of explored nodes S  $\leftarrow \emptyset$ 

  while (Q is not empty) {
    u  $\leftarrow$  delete min element from Q
    S  $\leftarrow S \cup \{u\}$ 
    foreach (edge e = (u, v) incident to u)
      if ((v  $\notin$  S) and ( $c_e < a[v]$ ))
        decrease priority a[v] to  $c_e$ 
  }
```

MST Algorithms: Theory

Deterministic comparison based algorithms.

- $O(m \log n)$ [Jarník, Prim, Dijkstra, Kruskal, Boruvka]
- $O(m \log \log n)$ [Cheriton-Tarjan 1976, Yao 1975]
- $O(m \beta(m, n))$ [Fredman-Tarjan 1987]
- $O(m \log \beta(m, n))$ [Gabow-Galil-Spencer-Tarjan 1986]
- $O(m \alpha(m, n))$ [Chazelle 2000]

Holy grail. $O(m)$.

Notable.

- $O(m)$ randomized. [Karger-Klein-Tarjan 1995]
- $O(m)$ verification. [Dixon-Rauch-Tarjan 1992]

Euclidean.

- 2-d: $O(n \log n)$. compute MST of edges in Delaunay
- k-d: $O(kn^2)$. dense Prim